Stefan Herytash

WU Vienna, Welthandelsplatz 1, 1020 Vienna, Austria

DRAFT

Generative AI as a Marketing Research Method

Course outline; the title is a working title. Master's elective in the cand.merc. EMF, next to *Brand Analytics*

Aim

Students learn to use language models as instruments of marketing research and to check them against human data. By the end of the course they can simulate respondents for an existing questionnaire, measure where the simulation departs from the human answers, use a model in the design and analysis of their own study, and correct a simulated estimate with a small human sample. The course grades whether students ran the check against human data.

Who it is for

Master's students in the cand.merc. EMF who know survey research and regression from their marketing research courses. The course sits one level above CBS's applied courses on AI in marketing and next to *Brand Analytics*, which already uses generative AI for market-research data.

Sessions

1. Simulated respondents against a real survey: one questionnaire, answered by a model and by the published sample
2. Where simulation and humans diverge: ranking against size, on 70 published marketing effects (r = .73, slope .38)
3. Model-assisted questionnaire design: items drafted with a model and pretested on simulated respondents
4. Analysis with language models: open answers coded by a model and checked against a human coder
5. Prediction-powered correction: simulated answers and a small human sample combined into an unbiased estimate
6. Agentic workflows: a model that plans and runs a study, and the points where the researcher checks it
7. Final project, defended orally: each team's design choices and the failure cases it tested

Assessment

Project	A team project: a study that uses at least one method of the course and checks it against human data.
Defence	An oral defence follows, in which each team explains its design choices and the failure cases it tested.
Grade	Students may use language models on their work, and the grade rests on the project and on how the team defends it.
Rubric	The rubric is handed out with the task, and each work step has a date.

Readings

Herytash, *When and How to Integrate AI Simulations into Behavioral Experimentation* (revise and resubmit, *International Journal of Research in Marketing*);

Collis, Eggers and Brynjolfsson (2024) and Eggers and Vriens (2026) on language-model valuations and priors in choice experiments;

Hewitt, Ashokkumar, Ghezae and Willer (2026) on predicting the results of survey experiments.

Tools

SimuLabs, the open-source engine I develop for simulated experiments; an open-weights language model, so that every run can be repeated; the survey files of published studies, which give each simulation a human benchmark.