Stefan Herytash

Institute for Strategy, Technology and Organization
WU Vienna, Welthandelsplatz 1, 1020 Vienna, Austria

October 2026

List of Publications

The list is numbered across all groups. The four entries marked with an asterisk (*) are submitted for review, in order of priority. Author position is stated on every entry.

Published and forthcoming

[1]* Koll, O., Datta, H., … Herytash, S., … van Heerde, H. J. (265 authors) (2026). *How Researchers' Analytical Decisions Impact Brand Price Elasticity Estimates: Insights from a Many-Analyst Study*. **International Journal of Research in Marketing**, forthcoming (Academic Journal Guide 4); page numbers not yet assigned. Preprint: doi:10.31234/osf.io/fs3k2_v4. Many-analyst study in which independent teams estimated brand price elasticities from the same data; I took part as one of the analysts. elasticity-open-science.com

Under review

[2]* Herytash, S. (2026). *When and How to Integrate AI Simulations into Behavioral Experimentation*. Solo-authored. **Revise & Resubmit at International Journal of Research in Marketing** (Academic Journal Guide 4). Asks when a language-model simulation can forecast the result of a behavioral experiment before it is run. Benchmarks simulated against human treatment effects from published studies, with each synthetic respondent completing the original questionnaire in one randomized condition; every statistic regenerates from one deterministic pipeline built on SimuLabs. Methods: cluster-robust meta-regression, publication-bias diagnostics, prediction-powered inference for a hybrid human–AI design.

[3]* Herytash, S. (2026). *After the Fax Machine Burns: Journal Certification When Everything Is Published*. Solo-authored. **Submitted to Management Science** (Financial Times 50; Academic Journal Guide 4*) as a Discussion Paper, in response to Bhargava et al. (2026) and Terwiesch (2026). Develops publish-everything-then-grade into a system a publisher could run: formal model of the five parties a publication system serves; Blackwell comparison of a public graded record against the accept verdict; settlement through prediction markets on replication outcomes. Methods: mechanism design, information economics.

[4]* Binz, M., Akata, E., … Herytash, S., … Schulz, E. (79 authors) (2026). *Post-training Makes Large Language Models Less Human-like*. Under review. arXiv:2605.07632. Introduces the Psych-201 dataset (208,021 participants, 25.9 million responses).

Working papers

[5] Herytash, S. (2026). *Threshold Migration to the Patient: Consumer Wearables, the Divided Gatekeeping Jurisdiction, and the Redistribution of Access*. Solo-authored. **Accepted at ANZAM 2026**, Health Management and Organisation stream, University of Newcastle. Consumer wearables split the clinical gatekeeping jurisdiction: an algorithm initiates triage while the profession retains disposition. Design: staggered difference-in-differences on international clearance dates for wearable atrial-fibrillation detection.

[6] Herytash, S. (2026). *Synthetic Participants Reproduce Effects They Cannot Replace*. Solo-authored working paper. A formal bound on how much a language-model simulation can add to a human sample: the effective-sample multiplier depends only on the within-study correlation between simulated and human responses, and it is finite whatever the simulator's size. Methods: prediction-powered inference, finite-sample results.

[7] Herytash, S. (2026). *The Matching Cost of Pay Transparency: Regulation, Firm Strategic Response and Labor Market Outcomes*. Solo-authored working paper. Formal model complete, empirical estimation in progress.

[8] Herytash, S. (2026). *SimuLabs: Auditable Simulation of Behavioral Experiments with Language Models, with Prediction-Powered Rectification*. Solo-authored. Prototype. Open-source engine (Rust) that runs studies from their original survey files and regenerates every reported statistic through one deterministic pipeline; the rectifier returns an unbiased estimate whatever the simulator does.

[9] Herytash, S. (2026). *The Platform Sets the Reference Wage: Salary Bands, Display Thresholds and Applicant Flows*. Solo-authored working paper. An employer-review platform displays a salary band only once enough employees have reported one, which makes the display threshold a discontinuity in what counts publicly as a normal wage. Methods: regression discontinuity at the display threshold.

Citation documentation

I do not keep a public Google Scholar profile. Citations are documented through my ORCID record, orcid.org/0000-0002-6642-4397, and through the DOI and arXiv identifiers given above.