Stefan Herytash

sherytash.github.io

15 October 2026

Research Profile, Publications and Pipeline

Research profile

I am an Assistant Professor at WU Vienna, trained in mathematics at TU Wien and the University of Zurich, with a PhD from LMU Munich's Institute for Market-Based Management, summa cum laude. My research measures when answers written by a language model can stand in for consumers in marketing research, as in a digital pre-test of an ad, an offer or a price. It also asks how platforms and journals certify quality once such evidence is cheap, from the salary band an employer-review platform displays to a journal's verdict on a manuscript.

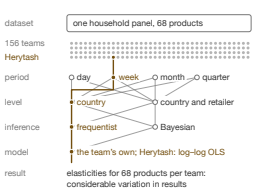
● **Measurement.** My paper on AI simulations in behavioral experiments tests how far simulated respondents can stand in for human participants in marketing experiments. I simulate 70 treatment effects from 30 published marketing studies with an open-weights language model, at the level of the instrument: each synthetic respondent completes the original questionnaire in one randomized condition. Simulated and human effects correlate at $r = .73$, but the calibration slope is .38, so the simulations understate effect sizes. They can rank ideas before a test is run; they cannot yet price them. A companion paper bounds what a simulation can add to a human sample: the effective-sample multiplier depends only on the within-study correlation between simulated and human responses, and it stays finite whatever the simulator's size. In the many-analyst study forthcoming in the **International Journal of Research in Marketing**, I was one of 156 analyst teams estimating brand price elasticities from the same household panel data. Across 10,415 estimates the team averages range from −7.3 to +3.1; choices on time aggregation, fixed effects, time controls and marketing-mix covariates explain the largest share of that spread.

■ **Certification.** kununu, the largest employer-review and salary platform in the German-speaking labour market, shows a salary band only once enough employees have reported one; its display threshold decides what counts publicly as a normal wage, and a regression discontinuity at that threshold can show what it does to applicant flows. *After the Fax Machine Burns* (submitted to **Management Science**) asks what a journal should certify when manuscripts cost almost nothing to write. It develops a system a publisher could run, in which every submission is published on arrival and graded afterward, and derives when each party a publication system serves prefers that graded record to the accept verdict. *Threshold Migration to the Patient* (accepted at ANZAM 2026) follows a consumer device into a profession's gate: a wearable's algorithm now initiates triage while the clinician keeps disposition.

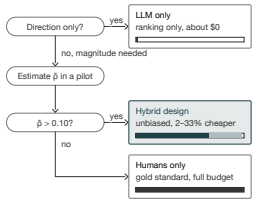
Publications and pipeline

● AI behavioral simulation ■ Post-AI certification

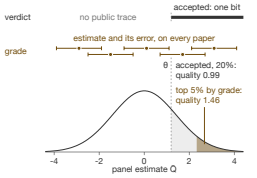
■ ***How Researchers' Analytical Decisions Impact Brand Price Elasticity Estimates: Insights from a Many-Analyst Study***
Koll, Datta, ..., Herytash, ..., van Heerde (265 authors); one of the analyst teams · Forthcoming, **International Journal of Research in Marketing** (Academic Journal Guide 4)



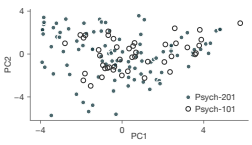
● ***When and How to Integrate AI Simulations into Behavioral Experimentation***
Revise & Resubmit at **International Journal of Research in Marketing** (Academic Journal Guide 4)



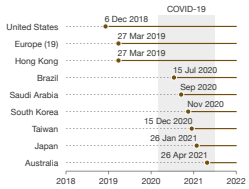
■ ***After the Fax Machine Burns: Journal Certification When Everything Is Published***
Submitted to **Management Science** (Financial Times 50; Academic Journal Guide 4*) as a Discussion Paper, in response to Bhargava et al. (2026) and Terwiesch (2026)



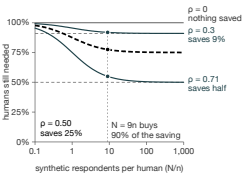
● *Post-training Makes Large Language Models Less Human-like*
Binz, Akata, …, Herytash, …, Schulz (79 authors) · Under Review; arXiv:2605.07632.
Introduces the Psych-201 dataset



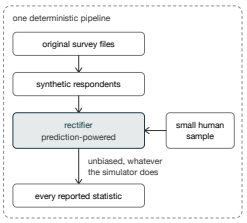
■ *Threshold Migration to the Patient: Consumer Wearables, the Divided Gatekeeping Jurisdiction, and the Redistribution of Access*
Accepted at ANZAM 2026, Health Management and Organisation stream, University of Newcastle



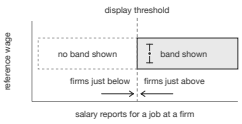
● *Synthetic Participants Reproduce Effects They Cannot Replace*
In Preparation



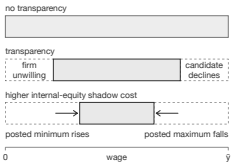
● *SimuLabs: Auditable Simulation of Behavioral Experiments with Language Models, with Prediction-Powered Rectification*
In Preparation; prototype



The Platform Sets the Reference Wage: Salary Bands, Display Thresholds and Applicant Flows
In Preparation; data partnership with kununu in negotiation



The Matching Cost of Pay Transparency: Regulation, Firm Strategic Response and Labor Market Outcomes
In Preparation; formal model complete, empirical estimation in progress with Revelio Labs data



A venue is named for each paper that is published, forthcoming, under review or accepted.

Intended research at CBS

My benchmark so far compares simulations with published results. With CogLab and the experiments.cbs.dk participant pool I would run the simulated and the human version of the same study side by side on Danish samples, starting with pre-tests of ads and offers. Promotion planning rests on price elasticities, so the elasticity study meets the Digital Marketing Research Cluster's project on online promotions and customer value; the kununu study and the Airbnb project below meet its project on the platform economy. In Denmark I would like to approach Salling Group, the country's largest retailer, for price elasticities, promotions and digital pre-tests of offers, and Carlsberg, a partner of the department's Consumer Research Cluster, for pre-tests of brand and pricing decisions with simulated and human respondents.

Beyond the pipeline I intend to study Airbnb's Superhost badge as a certification threshold on a digital platform; the project builds on my earlier work on Airbnb data at LMU. A host earns the badge with an average rating of at least 4.8 and with thresholds on response rate, cancellations and stays. The public Inside Airbnb snapshots carry each listing's rating and Superhost status, so a regression discontinuity or a bunching design at the thresholds can separate the badge from the quality it certifies. Trustpilot, the review platform founded in Denmark and based in Copenhagen, rates businesses from their customers' reviews; it is the Danish platform I would add to this work.

Another project would use Revelio Labs data on firms' job postings and employees' profiles, which WU licenses and the Matching Cost paper already uses. Zhu, Knight, Colicev and Bart (2026) find in Revelio

records of 164,780 firms that after the release of ChatGPT marketing headcount fell about 0.94 percent more than non-marketing headcount, mainly in non-digital marketing. I would ask which digital marketing tasks move to models and whether firms shift marketing roles toward AI skills.

The Matching Cost paper asks whether pay transparency changes whom firms hire, and Danish register data can test it. A Danish law of 2006 required firms with at least 35 employees to report wage statistics by gender, and Bennedsen, Simintzi, Tsoutsoura and Wolfenzon (2022, *Journal of Finance*) found with Statistics Denmark's matched employer-employee data that it narrowed the gender pay gap by two percentage points. Through CBS I could test the hiring question on these registers, next to the Revelio data the paper uses now.

Funding

My first Danish application would go to the Independent Research Fund Denmark's thematic call Research on Artificial Intelligence (up to DKK 5 million over five years, with a named co-applicant and a PhD student or postdoc), together with a senior colleague from the department. From 2029, when my PhD is three years old, the fund's Research Project 1 (up to DKK 2.55 million) and Sapere Aude: Research Leader (up to DKK 4.45 million) open to me; Sapere Aude is designed as preparation for a European Research Council Starting Grant. For applied work with Danish firms, Industriens Fond accepts consortium applications on a rolling basis.

Open science

SimuLabs, the engine behind the simulation papers, will be made open source. I will publish the data and code of each paper when it is published.