Stefan Herytash

sherytash.github.io

15 October 2026

Research Profile, Publications and Pipeline

Research profile

A response to a questionnaire used to cost a participant's time, and a manuscript cost its author months. Part of what made either informative was that cost. Language models have made both cheap. My research follows two consequences for marketing: what can be measured now that a costly signal is cheap, and how the institutions that certify evidence must be rebuilt.

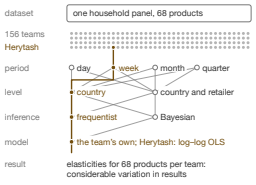
● **Measurement.** My paper on AI simulations in behavioral experiments tests how far simulated respondents can stand in for human participants in behavioural experiments, and where that substitution stops. I simulate 70 treatment effects from 30 published marketing studies with an open-weights language model, at the level of the instrument: each synthetic respondent completes the original questionnaire in one randomized condition. Simulated and human effects correlate at $r = .73$, but the calibration slope is .38, so the simulations understate effect sizes. They can rank ideas before a test is run; they cannot yet price them. A companion paper bounds what a simulation can add to a human sample: the effective-sample multiplier depends only on the within-study correlation between simulated and human responses, and it stays finite whatever the simulator's size. The econometric side of measurement is the many-analyst study forthcoming in the *International Journal of Research in Marketing*, in which I was one of 156 analyst teams estimating brand price elasticities from the same household panel data. Across 10,415 estimates the team averages range from −7.3 to +3.1; choices on time aggregation, fixed effects, time controls and marketing-mix covariates explain the largest share of that spread.

■ **Certification.** When manuscripts cost almost nothing to write, a journal's accept verdict carries less of the information it used to. *After the Fax Machine Burns* develops a system a publisher could run, in which every submission is published on arrival and graded afterward, and derives when each party a publication system serves prefers that graded record to the accept verdict. Digital platforms certify in the same way. kununu, the largest employer-review and salary platform in the German-speaking labour market, shows a salary band only once enough employees have reported one; its display threshold decides what counts publicly as a normal wage, and a regression discontinuity at that threshold can show what it does to applicant flows. *Threshold Migration to the Patient* follows a consumer device into a profession's gate: a wearable's algorithm now initiates triage while the clinician keeps disposition.

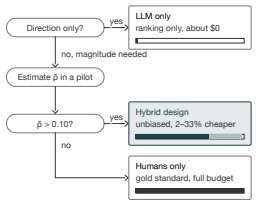
Publications and pipeline

● AI behavioral simulation ■ Post-AI certification

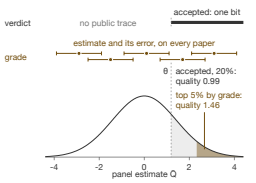
■ ***How Researchers' Analytical Decisions Impact Brand Price Elasticity Estimates: Insights from a Many-Analyst Study***
Koll, Datta, ..., Herytash, ..., van Heerde (265 authors); one of the analyst teams · Forthcoming, *International Journal of Research in Marketing* (Academic Journal Guide 4)



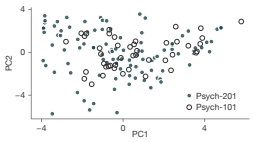
● ***When and How to Integrate AI Simulations into Behavioral Experimentation***
Solo-authored · Revise & Resubmit at International Journal of Research in Marketing (Academic Journal Guide 4)



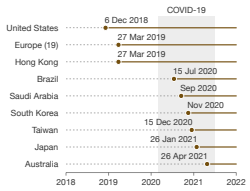
■ ***After the Fax Machine Burns: Journal Certification When Everything Is Published***
Solo-authored · Submitted to Management Science (Financial Times 50; Academic Journal Guide 4*) as a Discussion Paper, in response to Bhargava et al. (2026) and Terwiesch (2026)



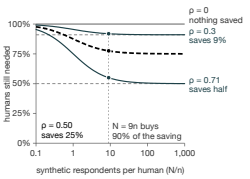
● *Post-training Makes Large Language Models Less Human-like*
Binz, Akata, ..., Herytash, ..., Schulz (79 authors) · Under Review; arXiv:2605.07632.
Introduces the Psych-201 dataset



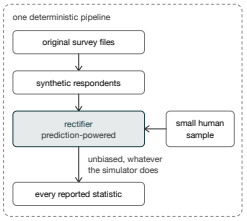
■ *Threshold Migration to the Patient: Consumer Wearables, the Divided Gatekeeping Jurisdiction, and the Redistribution of Access*
Solo-authored · Accepted at ANZAM 2026, Health Management and Organisation stream, University of Newcastle



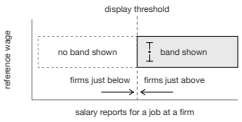
● *Synthetic Participants Reproduce Effects They Cannot Replace*
Solo-authored · In Preparation



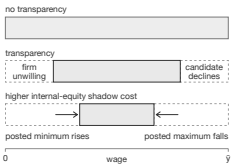
● *SimuLabs: Auditable Simulation of Behavioral Experiments with Language Models, with Prediction-Powered Rectification*
Solo-authored · In Preparation; prototype



The Platform Sets the Reference Wage: Salary Bands, Display Thresholds and Applicant Flows
Solo-authored · In Preparation; data partnership with kununu in negotiation



The Matching Cost of Pay Transparency: Regulation, Firm Strategic Response and Labor Market Outcomes
Solo-authored · In Preparation; formal model complete, empirical estimation in progress with Revelio Labs data



A venue is named for each paper that is published, forthcoming, under review or accepted.

Intended research at CBS

My work meets the Digital Marketing Research Cluster at two points. Felix Eggers asks whether a language model can stand in for a human respondent when valuing digital goods; my benchmark asks the same question for behavioural experiments and measures where the substitution holds. Antonia Erz studies how two-sided review platforms change behaviour; the employer-review platform in my salary-transparency project is a market of that kind. Where the department tests the validity of new measurement, as Constant Pieters does, simulated respondents are a case that still has to pass that test. The elasticity work adds a strand the Digital Marketing Research Cluster lists among its projects, online promotions and customer value, since promotion planning rests on price elasticities.

Beyond the pipeline I intend to open two projects in digital marketing. The first treats Airbnb's Superhost badge as a certification threshold on a digital platform. A host earns the badge with an average rating of at least 4.8 and with thresholds on response rate, cancellations and stays. The public Inside Airbnb snapshots carry each listing's rating and Superhost status, so a regression discontinuity or a bunching design at the thresholds can separate the badge from the quality it certifies. The second uses Revelio Labs data on firms' job postings and employees' profiles, which WU licenses through WRDS. It asks whether firms shift marketing roles toward AI skills after generative AI, and which digital marketing tasks move to models; it would sit next to Zhu, Knight, Colicev and Bart (2026) on AI and marketing employment. The Superhost project builds on my earlier work on Airbnb data at LMU, and the Revelio project on the Matching Cost paper, which already uses Revelio Labs data. Also at LMU, I worked on predicting engagement with social media ads using machine learning.

Funding

My first Danish application would go to the Independent Research Fund Denmark's thematic call Research on Artificial Intelligence, which funds projects of up to DKK 5 million over up to five years with a named co-PI and at least one PhD student or postdoc; I would apply with a senior colleague from the department. Once my PhD is three years old, from 2029, DFF Research Project 1 (up to DKK 2.55 million) and Sapere Aude: DFF-Research Leader (up to DKK 4.45 million) become open to me, and Sapere Aude is designed as preparation for an ERC Starting Grant. For applied work with Danish firms, Industriens Fond accepts consortium applications on a rolling basis. At LMU I raised over €100,000 in research funding for the institute from national and international industry partners.

Open science

SimuLabs, the engine behind the simulation papers, is open source (Rust). It runs studies from their original survey files, and every statistic reported in that paper regenerates from one deterministic pipeline. Its prediction-powered rectifier returns an unbiased estimate whatever the simulator does. Each paper's replication package is released on publication. The elasticity study is itself an open-science design, with all teams working from the same data, and the Psych-201 dataset is public.